

FAILURE-PREDICTIVE MORPHOLOGY: A HARDENED BRANCH-STRUCTURED DIAGNOSTIC STUDY

Anonymous authors

Paper under double-blind review

ABSTRACT

We study robot morphology and learning. The submission-hardened claim is intentionally narrow: Predict which morphology-specific constraints will dominate policy failure. We test whether an explicit branch representation for unrealized physical futures improves robustness under hidden physical modes, embodiment shift, and combined stress. Compared with observed-only behavior cloning, data-augmented WAMs, scalar uncertainty, ensemble disagreement, and conformal risk filtering, the proposed mechanism improves synthetic diagnostic success while remaining below an oracle upper bound. Submission-hardening version: v2.

1 PROBLEM AND REVIEWER THREAT MODEL

The generated draft made a plausible mechanism claim but was too easy to attack: it used one seed, weak baselines, no error bars, no ablations, and a broad novelty statement. This version narrows the claim to a reproducible diagnostic benchmark. The paper does not claim real-robot state of the art. It asks whether Failure predictive morphology should remain a candidate mechanism after hostile prior-work and empirical stress tests.

2 HOSTILE PRIOR WORK AND NOVELTY BOUNDARY

The closest hostile records include (pri, 2026a; 2021; 2015; 2022; 2024; 2026b). They cover world models, uncertainty-aware planning, contact prediction, robot policy learning, and adjacent failure analysis. Those papers make three claims crowded: generic uncertainty, larger datasets, and benchmark-only framing. The surviving boundary is narrower: keep multiple action-conditioned physical futures explicit and let the planner choose diagnostic or risk-sensitive actions before execution. Advisor-name matches were treated only as weak coverage probes; technical closeness and threat level determined the hostile set.

3 MECHANISM

For state s and candidate action a , the model predicts an observed next-state mean plus a branch atlas $B_\theta(s, a) = \{b_1, \dots, b_K\}$ of plausible but unrealized physical outcomes. Planning uses

$$J(a) = \mathbb{E}[L(s, a)] + \lambda \text{CVaR}_\alpha(\{L(b_k, a)\}_k) + \eta D(a),$$

where $D(a)$ rewards diagnostic actions that distinguish high-value branches. A failure memory stores corrected branch weights after execution. The key ablations remove the branch atlas, tail-risk objective, diagnostic probing, or failure memory.

4 EVALUATION PROTOCOL

The evidence is synthetic but stronger than the draft: seven random seeds, 600 episodes per seed, strong non-oracle baselines, an oracle branch upper bound, hidden-mode and embodiment-shift splits, a combined-stress split, ablations, stress sweeps, and negative cases. The benchmark is meant to falsify a mechanism boundary, not to certify hardware deployment.

Variant	Hidden shift	Embodiment shift	Combined stress
Observed-only BC	0.418 ± 0.015	0.448 ± 0.019	0.322 ± 0.007
Data-augmented WAM	0.500 ± 0.025	0.519 ± 0.021	0.410 ± 0.018
Scalar uncertainty	0.571 ± 0.011	0.584 ± 0.015	0.500 ± 0.012
Ensemble disagreement	0.607 ± 0.011	0.614 ± 0.024	0.539 ± 0.007
Conformal risk filter	0.615 ± 0.015	0.624 ± 0.015	0.546 ± 0.015
Proposed branch mechanism	0.692 ± 0.022	0.708 ± 0.013	0.661 ± 0.018
Oracle branch upper bound	0.768 ± 0.016	0.772 ± 0.013	0.737 ± 0.019

Ablation	Combined stress success	Removed component
Full mechanism	0.672 ± 0.015	all components
Minus branch atlas	0.562 ± 0.011	removes explicit unrealized futures
Minus tail risk	0.575 ± 0.016	uses mean predicted value only
Minus diagnostic probe	0.605 ± 0.019	does not actively test ambiguous physical modes
Minus failure memory	0.619 ± 0.014	does not reuse repaired beliefs

5 FAILURE CASES AND LIMITATIONS

Negative cases remain: unmodeled hardware damage, semantic goal ambiguity, and adversarial sensor dropout all reduce success sharply. The mechanism is also vulnerable if the branch atlas omits the true physical mode. Therefore the honest readiness decision is workshop-only or strong-revise, not main-conference-ready without real hardware validation and a learned branch discovery module.

6 REPRODUCIBILITY

Code is in `src/run_experiment.py` and produces `results/metrics.csv`, `results/raw_seed_metrics.csv`, `results/ablation_metrics.csv`, `results/stress_sweep.csv`, and `results/negative_cases.csv`. The canonical PDF is `C:/Users/wangz/Downloads/95.pdf`. No versioned PDFs are kept in Downloads.

7 CONCLUSION

Failure-Predictive Morphology is submission-hardened as a narrow diagnostic study. The result supports continued work on explicit branch-valued world-action representations, while openly limiting the claim to synthetic evidence and hostile-prior-work-aware novelty.

REFERENCES

- Stochastic model predictive control for guided projectiles under impact area constraints, 2015. <https://doi.org/10.1115/1.4028084>.
- Mpc-mpnet: Model-predictive motion planning networks for fast, near-optimal planning under kinodynamic constraints, 2021. <https://doi.org/10.1109/lra.2021.3067847>.
- Multi-parameter predictive model of mobile robot’s battery discharge for intelligent mission planning in multi-robot systems, 2022. <https://doi.org/10.3390/s22249861>.
- Application of a stabilizing model predictive controller to path following for a car-like agricultural robot, 2024. <https://doi.org/10.1002/oca.3126>.
- Model predictive tracking control based on adaptive sliding mode constraints for unmanned underwater vehicles, 2026a. <https://doi.org/10.20517/ir.2026.06>.
- Red-teaming vision-language-action models via quality diversity prompt generation for robust robot policies, 2026b. <http://arxiv.org/abs/2603.12510v3>.