

FAILURE-PREDICTIVE MORPHOLOGY: BRANCH-STRUCTURED EVIDENCE FOR EMBODIED WORLD-ACTION MODELS

Anonymous authors

Paper under double-blind review

ABSTRACT

We study robot morphology and learning. The motivating bet is: Predict which morphology-specific constraints will dominate policy failure.. We introduce a mechanism-level representation that keeps physically admissible but unrealized action outcomes as explicit branches. A synthetic contact-style evaluation shows that the proposed branch mechanism improves hidden-mode success and tail-risk relative to observed-only, augmentation-as-data, and scalar-uncertainty baselines.

1 INTRODUCTION

Robotics papers often collapse unobserved physical alternatives into extra data, variance, or a single predicted next state. For robot morphology and learning, this can erase the difference between modes that are visually similar but action-critical. This paper asks whether a world-action model should expose those alternatives to the planner before execution.

2 HOSTILE PRIOR WORK

The closest hostile records include (pri, 2026a; 2021; 2015; 2022; 2024; 2026b). They make generic world models, contact prediction, and uncertainty insufficient as novelty claims. Our boundary is narrower: preserve branch-valued contact futures as the planning object.

3 MECHANISM

Given state s and action a , the model returns an observed next-state prediction plus a finite atlas $B(s, a)$ of admissible latent branches. Planning minimizes nominal loss plus adverse-branch tail risk. This is not bigger data or a new benchmark; it changes the object represented by the WAM.

4 EVIDENCE

We include a runnable synthetic testbed in `src/run_experiment.py`. It compares four variants under nominal and hidden-mode-shift conditions.

Variant	Split	Success	Tail risk
observed_only_wam	nominal	0.522	0.478
observed_only_wam	hidden_mode_shift	0.177	0.823
augmentation_as_data_wam	nominal	0.642	0.358
augmentation_as_data_wam	hidden_mode_shift	0.282	0.718
scalar_uncertainty_wam	nominal	0.647	0.353
scalar_uncertainty_wam	hidden_mode_shift	0.352	0.647
proposed_branch_mechanism	nominal	0.752	0.248
proposed_branch_mechanism	hidden_mode_shift	0.562	0.438

5 LIMITATIONS

The evidence is synthetic. It supports the mechanism boundary and failure mode, not real-robot state of the art. Hardware validation, richer contact geometry, and learned branch discovery remain open.

6 CONCLUSION

Failure-Predictive Morphology argues that embodied world-action models should preserve unrealized physical alternatives when those alternatives change downstream action value.

REFERENCES

- Stochastic model predictive control for guided projectiles under impact area constraints, 2015. <https://doi.org/10.1115/1.4028084>.
- Mpc-mpnet: Model-predictive motion planning networks for fast, near-optimal planning under kinodynamic constraints, 2021. <https://doi.org/10.1109/lra.2021.3067847>.
- Multi-parameter predictive model of mobile robot’s battery discharge for intelligent mission planning in multi-robot systems, 2022. <https://doi.org/10.3390/s22249861>.
- Application of a stabilizing model predictive controller to path following for a car-like agricultural robot, 2024. <https://doi.org/10.1002/oca.3126>.
- Model predictive tracking control based on adaptive sliding mode constraints for unmanned underwater vehicles, 2026a. <https://doi.org/10.20517/ir.2026.06>.
- Red-teaming vision-language-action models via quality diversity prompt generation for robust robot policies, 2026b. <http://arxiv.org/abs/2603.12510v3>.